

Explore Hugging Face

Estimated time: 30 minutes

Welcome to the hands-on lab, Explore Hugging Face.

Learning objectives

After completing this lab, you will be able to:

- Explore the features and capabilities of Hugging Face, an AI community and data science platform.

Introduction

Generative artificial intelligence (AI) is on the rise because of its ability to train models with extensive data. Generative AI excels at in-context text generation, chatbot development, text summarization, marketing content creation, and more, spanning language and media domains. Yet, you would require a platform to make the most of these models. This platform should serve data scientists looking to access models via application programming interfaces (APIs), application developers seeking to leverage these capabilities, researchers exploring model functionalities, and contributors looking to support ongoing development. A community platform is vital for accessing, building, training, deploying, reusing, and contributing to these generative AI innovations.

In this lab, you'll be introduced to one such AI community platform, Hugging Face, and its user interface (UI).

Hugging Face


Hugging Face is an AI community and data science platform where data scientists, machine learning engineers, and researchers can collaborate to build, train, and deploy models, datasets, and applications. The platform also welcomes those who wish to contribute or seek support for ongoing development.

Exercise 1: Explore the Hugging Face platform and its capabilities

Hugging Face provides the framework to demo, run, and deploy AI in real-world applications. It also allows developers to openly share and test their models while providing access to models and datasets contributed by other users.

Step 1: Getting started

1. Click [Hugging Face - The AI community building the future****](#) to launch the Hugging Face platform.
Note: Although you can use the platform without logging in, we recommend logging in for the best experience.
2. If you are a registered user, click **Log In** in the top right-hand corner.



3. Then, on the **Log In** page, enter your **username** or registered **email address** and **password**. Then, click **Login**.




Log In

Don't have an account? [Sign Up](#)

[Forgot your password?](#)

[SSO](#) is available for companies

4. If you are a first-time user, click **Sign Up** on the top right-hand corner of the Hugging Face landing page.



The Hugging Face landing page features a dark theme. At the top, there is a navigation bar with the Hugging Face logo, a search bar, and links to Models, Datasets, Spaces, Docs, Solutions, Pricing, Log In, and a highlighted Sign Up button. The main content area is divided into two sections. On the left, a large banner reads "The AI community building the future." and "The platform where the machine learning community collaborates on models, datasets, and applications." On the right, there is a list of featured models, including meta-llama/llama-2-70b, stabilityai/stable-diffusion-xl-base-0.9, openchat/openchat, llllyasviel/ControlNet-v1-1, cerspense/zeroscope_v2_XL, meta-llama/llama-2-13b, tiuae/falcon-40b-instruct, WizardLM/WizardCoder-150-V1.0, CompVis/stable-diffusion-v1-4, and stabilityai/stable-diffusion-2-1. A sidebar on the left lists various tasks and categories like Multimodal, Computer Vision, Natural Language Processing, and Audio.

5. On the **Join Hugging Face** page, enter your email address and a desired password. Then, click **Next**.
Note: The password must contain uppercase letters, lowercase letters, and numbers.



Join Hugging Face

Join the community of machine learners!

Email Address

Hint: Use your organization email to easily find and join your company/team org.

Password

Next

Already have an account? [Log in](#)
[SSO](#) is available for companies

6. On the **Complete your profile** page, enter a **username** and your **full name**, which are mandatory fields. Then, select the **I have read and agree with the Terms of Service and the Code of Conduct** checkbox and click **Create Account**.

Complete your profile

One last step to join the community

Username

Full name

Avatar (optional)

Upload file

GitHub username (optional)

Homepage (optional)


Twitter username (optional)

Research interests (optional)


☐ I have read and agree with the [Terms of Service](#) and the [Code of Conduct](#)

Create Account

7. Verify you're a human by completing the hCaptcha challenge.




8. On successful login, you will be taken to your homepage.




Step 2: Getting acquainted with the Hugging Face UI


1. At the top of the homepage, on the menu bar, you'll find the search bar, the various features available on Hugging Face, and your avatar. The search functionality lets you quickly find models, datasets, and spaces.




2. In the left pane, you'll see your profile, organization profile, if you've created one, and the available resources.



3. In the central portion of the homepage, your activity is displayed, including your models, datasets, spaces, and community discussions.



4. In the right pane, you can view the trending models, datasets, and spaces.



Let's explore the six features: Models, Databases, Spaces, Docs, Solutions, and Pricing in detail.

Step 3: Exploring Models

Hugging Face hosts several models designed for a range of machine learning (ML) tasks. The Models page has two panes. On the left pane, you'll find six quick access tabs: Tasks, Libraries, Datasets, Languages, Licenses, and Other. Each tab has a filter to help you quickly find what you want. The right pane displays the list of all available models and a filter bar.

The screenshot shows the Hugging Face website's 'Models' page. The top navigation bar includes 'Models', 'Datasets', 'Spaces', 'Docs', 'Solutions', and 'Pricing'. The left sidebar has tabs for 'Tasks', 'Libraries', 'Datasets', 'Languages', 'Licenses', and 'Other'. The 'Tasks' tab is active, and a search filter 'summarization' is applied. The main content area displays a grid of model cards. The first card in the first column is 'facebook/bart-large-cnn', which is highlighted with a red box. Other visible model cards include 'HuggingFaceH4/zephyr-7b-alpha', 'mistralai/Mistral-7B-v0.1', and 'mistralai/Mistral-7B-Instruct-v0.1'.

In this lab, let's explore the summarization models available.

1. In the top navigation, click **Models**.
2. On the **Models** main page, in the left pane, click the **Tasks** tab, and then type "summarization" in the filter bar.
3. From the filtered results, click **Summarization** under the **Natural Language Processing** section.
4. The right pane will display all the models capable of text summarization.
5. From the displayed list of models, click **facebook/bart-large-cnn**.

This screenshot shows the detailed page for the 'facebook/bart-large-cnn' model. The left sidebar shows the 'Tasks' tab with 'summarization' entered in the filter. Under the 'Natural Language Processing' section, 'Summarization' is selected. The main area displays the model card for 'facebook/bart-large-cnn', which is highlighted with a red box. The card shows the model's name, 'Summarization' task, and statistics like 'Updated 16 days ago' and '1.87M' likes. Other model cards are visible in the background, such as 'pszemraj/led-base-book-summary' and 'google/pegasus-large'.

6. On the selected model's page, you'll find three tabs.

- The first tab, **Model card**, will provide a model description, uses and limitations, how-to-use guidelines, BibTeX entry and citation info, and the hosted inference API pane, which is the text summarization field.
- The second tab, **Files and versions**, will display all data about model commits.
- The third tab, **Community**, will display all community discussions and pull requests.

Hugging Face Search models, datasets, users...

facebook/bart-large-cnn like 618

Summarization Transformers PyTorch TensorFlow JAX Rust cnn_dailymail English bart text2text-generation Eval Results Inference Endpoints arxiv:1910.13461

License: mit

Model card Files and versions Community 57

Edit model card

BART (large-sized model), fine-tuned on CNN Daily Mail

BART model pre-trained on English language, and fine-tuned on [CNN Daily Mail](#). It was introduced in the paper [BART: Denoising Sequence-to-Sequence Pre-training for Natural Language Generation, Translation, and Comprehension](#) by Lewis et al. and first released in [this repository] (<https://github.com/pytorch/fairseq/tree/master/examples/bart>).

Disclaimer: The team releasing BART did not write a model card for this model so this model card has been written by the Hugging Face team.

Model description

BART is a transformer encoder-encoder (seq2seq) model with a bidirectional (BERT-like) encoder and an autoregressive (GPT-like) decoder. BART is pre-trained by (1) corrupting text with an arbitrary noising function, and (2) learning a model to reconstruct the original text.

BART is particularly effective when fine-tuned for text generation (e.g. summarization, translation) but also works well for comprehension tasks (e.g. text classification, question answering). This

Downloads last month
1,873,269

Hosted inference API

Summarization Examples

The tower is 324 metres (1,063 ft) tall, about the same height as an 81-storey building, and the tallest structure in Paris. Its base is square, measuring 125 metres (410 ft) on each side. During its construction, the Eiffel Tower surpassed the Washington Monument to become the tallest man-made structure in the world, a title it held for 41 years until the Chrysler Building in New York City was finished in 1930. It was the first structure to reach a height of 300 metres. Due to the addition of a broadcasting aerial at the top of the tower in 1957, it is now taller than the Chrysler Building by 5.2 metres (17 ft). Excluding transmitters, the Eiffel Tower is the second tallest free-standing structure in France after the Millau Viaduct.

Compute

Computation time on Intel Xeon 3rd Gen Scalable cpu: cached

7. To try the model's text summarization capability, in the **Hosted inference API** section, delete the default text in the **Summarization** field.

Model card Files and versions Community 57

Edit model card

BART (large-sized model), fine-tuned on CNN Daily Mail

BART model pre-trained on English language, and fine-tuned on [CNN Daily Mail](#). It was introduced in the paper [BART: Denoising Sequence-to-Sequence Pre-training for Natural Language Generation, Translation, and Comprehension](#) by Lewis et al. and first released in [this repository] (<https://github.com/pytorch/fairseq/tree/master/examples/bart>).

Disclaimer: The team releasing BART did not write a model card for this model so this model card has been written by the Hugging Face team.

Model description

BART is a transformer encoder-encoder (seq2seq) model with a bidirectional (BERT-like) encoder and an autoregressive (GPT-like) decoder. BART is pre-trained by (1) corrupting text with an arbitrary noising function, and (2) learning a model to reconstruct the original text.

BART is particularly effective when fine-tuned for text generation (e.g. summarization, translation) but also works well for comprehension tasks (e.g. text classification, question answering). This particular checkpoint has been fine-tuned on CNN Daily Mail, a large collection of text-summary pairs.

Intended uses & limitations

Downloads last month
1,873,269

Hosted inference API

Summarization Examples

Your sentence here...

Compute

Computation time on Intel Xeon 3rd Gen Scalable cpu: cached

The tower is 324 metres (1,063 ft) tall, about the same height as an 81-storey building. Its base is square, measuring 125 metres (410 ft) on each side. During its construction, the Eiffel Tower surpassed the Washington Monument to become the tallest man-made structure in the world.

JSON Output

Dataset used to train facebook/bart-large-cnn

cnn_dailymail

Viewer · Updated Nov 18, 2022 · 97.5k · 114

8. Then, enter the text you want to summarize in the **Summarization** field and click **Compute**.
Note: The summarized text will take some time to generate.

BART (large-sized model), fine-tuned on CNN Daily Mail

BART model pre-trained on English language, and fine-tuned on [CNN Daily Mail](#). It was introduced in the paper [BART: Denoising Sequence-to-Sequence Pre-training for Natural Language Generation, Translation, and Comprehension](#) by Lewis et al. and first released in [this repository (<https://github.com/pytorch/fairseq/tree/master/examples/bart>)].

Disclaimer: The team releasing BART did not write a model card for this model so this model card has been written by the Hugging Face team.

Model description

BART is a transformer encoder-encoder (seq2seq) model with a bidirectional (BERT-like) encoder and an autoregressive (GPT-like) decoder. BART is pre-trained by (1) corrupting text with an arbitrary noising function, and (2) learning a model to reconstruct the original text.

BART is particularly effective when fine-tuned for text generation (e.g. summarization, translation) but also works well for comprehension tasks (e.g. text classification, question answering). This particular checkpoint has been fine-tuned on CNN Daily Mail, a large collection of text-summary pairs.

Intended uses & limitations

You can use this model for text summarization.

How to use

Downloads last month
1,873,269



Hosted inference API

Summarization

Examples

The Eiffel Tower, one of the most iconic landmarks in the world, stands as a symbol of France's artistic and engineering prowess. Located in the heart of Paris, this towering masterpiece of iron architecture has captivated the world since its completion in 1889. This essay will delve into the history, design, significance, and enduring allure of the Eiffel Tower. The Eiffel Tower was designed by the renowned engineer Gustave Eiffel and was constructed as the centerpiece of the 1889 Exposition Universelle, a world fair held in Paris to commemorate the 100th anniversary of the French Revolution. It was initially met with mixed reactions, with some prominent figures expressing disdain for its unconventional design. However, Eiffel's vision prevailed, and the tower was completed in just over two years, an astonishing feat given the technology of the late 19th century. The Eiffel Tower is a true marvel of structural engineering. Standing at a height of 324 meters (1,063 feet), it was the tallest man-made structure in the world until the completion of the Empire State Building in 1931. The tower is constructed primarily of iron, making it a testament to the possibilities of this once-underestimated material. Its lattice-like structure features over 18,000 individual iron pieces, all seamlessly connected to form a graceful and airy silhouette. The tower's design is divided into three levels, each offering a unique perspective of Paris. The first level features restaurants, shops, and a glass floor, providing visitors with a surreal experience of looking down at the city. The second level offers stunning panoramic views of Paris and is home

Compute

Computation time on Intel Xeon 3rd Gen Scalable cpu: 6.990 s

Hosted inference API

Summarization

The Eiffel Tower, one of the most iconic landmarks in the world, stand symbol of France's artistic and engineering prowess. Located in the he Paris, this towering masterpiece of iron architecture has captivated th since its completion in 1889. This essay will delve into the history, desi significance, and enduring allure of the Eiffel Tower. The Eiffel Tower v designed by the renowned engineer Gustave Eiffel and was constructe centerpiece of the 1889 Exposition Universelle, a world fair held in Par commemorate the 100th anniversary of the French Revolution. It was met with mixed reactions, with some prominent figures expressing dis unconventional design. However, Eiffel's vision prevailed, and the tow completed in just over two years, an astonishing feat given the techno late 19th century. The Eiffel Tower is a true marvel of structural engine Standing at a height of 324 meters (1,063 feet), it was the tallest man-r structure in the world until the completion of the Empire State Buildin The tower is constructed primarily of iron, making it a testament to th possibilities of this once-underestimated material. Its lattice-like struc features over 18,000 individual iron pieces, all seamlessly connected to graceful and airy silhouette. The tower's design is divided into three le offering a unique perspective of Paris. The first level features restaurar and a glass floor, providing visitors with a surreal experience of lookin

Compute

Computation time on Intel Xeon 3rd Gen Scalable cpu: 13.222 s

The Eiffel Tower was the tallest man-made structure in the world until completion of the Empire State Building in 1931. Its lattice-like structu over 18,000 individual iron pieces, all seamlessly connected to form a airy silhouette. The tower's design is divided into three levels, each off perspective of Paris. Today, the tower is a major tourist attraction draw of visitors from all corners of the globe.

9. The condensed text will be generated in the box below the summarization field.

the paper [BART: Denoising Sequence-to-Sequence Pre-training for Natural Language Generation, Translation, and Comprehension](#) by Lewis et al. and first released in [this repository (<https://github.com/pytorch/fairseq/tree/master/examples/bart>)].

Disclaimer: The team releasing BART did not write a model card for this model so this model card has been written by the Hugging Face team.

Model description

BART is a transformer encoder-encoder (seq2seq) model with a bidirectional (BERT-like) encoder and an autoregressive (GPT-like) decoder. BART is pre-trained by (1) corrupting text with an arbitrary noising function, and (2) learning a model to reconstruct the original text.

BART is particularly effective when fine-tuned for text generation (e.g. summarization, translation) but also works well for comprehension tasks (e.g. text classification, question answering). This particular checkpoint has been fine-tuned on CNN Daily Mail, a large collection of text-summary pairs.

Intended uses & limitations

You can use this model for text summarization.

How to use

Here is how to use this model with the [pipeline API](#):

```
from transformers import pipeline

summarizer = pipeline("summarization", model="facebook/bart-large-cnn")
```

Step 4: Try yourself

Now, try exploring other text summarization models. You could also try other models.

For example,


Summarization: [google/pegasus-large](#)

Text-to-image: [stabilityai/stable-diffusion-xl-base-1.0](#)

Step 5: Exploring Databases

The Hugging Face platform hosts many datasets curated by community members for various tasks like automatic speech recognition, translation, and image classification. In addition to the dataset card details, many datasets offer a Dataset Preview to showcase the data. Each dataset is stored as a Git repository, complete with the essential scripts for data download and the creation of splits for training, evaluation, and testing purposes.

Like the Models page, the Datasets page is organized into two panes. On the left pane, you'll find six quick access tabs: Tasks, Sizes, Sub-tasks, Languages, Licenses, and Other, each equipped with a filter for efficient searching. The right pane displays the list of all available datasets and a filter bar.



1. In the top navigation, click **Datasets**.
2. On the **Datasets** main page, in the left pane, click the **Tasks** tab, and then type “**summarization**” in the filter bar.
3. From the filtered results, click **Summarization** under the **Natural Language Processing** section.
4. The right pane will display all the datasets used to train models for text summarization.
5. From the displayed list of datasets, select **Open-Orca/OpenOrca**.

6. On the selected Dataset's page, you'll find two tabs.

- On the first tab, **Dataset card**, you can view the dataset viewer, official models, dataset attribution, supported tasks and leaderboards, languages, dataset structure, dataset use, citation, and so on.
- The second tab, **Files and versions**, will display the history of commits.

Step 6: Try yourself

Now, take a look at other datasets used to train models.

For example,

Text-to-Video: *shinonomelab/cleanvid-15m_map*

Image-to-Text: *MMinstruction/M3IT*

Step 7: Exploring Spaces

Hugging Face Spaces is the quickest and easiest way to build, deploy, and share ML demo applications on your personal or organization's profiles. It's also a versatile option for community members to share projects freely. The best part is users with no prior technical knowledge can easily create and showcase applications, allowing other users to test them.

The Spaces main page displays a search bar and all available spaces.


Hugging Face

[Models](#)
[Datasets](#)
[Spaces](#)
[Docs](#)
[Solutions](#)
[Pricing](#)



Spaces

Discover amazing ML apps made by the community!

[Create new Space](#)
or [learn more](#)

[new](#)
[Full-text search](#)

11
5

☆ Spaces of the week


Running on T4
80



artificialguybr
11 days ago

Building on A10G
49



ChenyangSi
4 minutes ago

Running on A100
186



HuggingFaceH4
8 days ago

Running on A10G



artificialguybr
abou

Running on T4
61



mikonvergence
10 days ago

Running on A10G
85



chansung
1 day ago

Running on T4
107



badayvedat
8 days ago

All running apps, trending first

Note: The application status (running, stopped, configuration error, runtime error, and so on) appears on the top-left of each application card, and the creator's name appears below each card. Please make sure to choose an application that is currently running.

1. Click <https://huggingface.co/spaces/Amit29sonawane/stabilityai-stable-diffusion-xl-base-1.0>.

Note: This text-to-image generation application has been specifically created on Spaces for this lab.

2. In the **Input** field, enter a prompt and click **Submit**.

For example, type the prompt **"Mickey Mouse waving."**

3. The image will be generated in the **Output** space.

Note: The image will take some time to generate.

Input

Mickey Mouse waving

Clear

Submit



Step 8: Try yourself

Now, try generating different images on this application. You can also test different applications on Spaces.

For example,

Text-to-image: <https://huggingface.co/spaces/Amit29sonawane/stabilityai-stable-diffusion-xl-base-1.0>

Text summarizer: <https://huggingface.co/spaces/Amit29sonawane/textsummarizer>

Note: These applications have been specifically created on Spaces for this lab.

Step 9: Exploring Docs

Hugging Face Docs offers a user-friendly, in-depth resource, structuring information into several sections:

- **Get Started:** Provides a concise library overview and installation instructions
- **Tutorials:** Provides common use cases as examples
- **How-to-Guides:** Furnishes detailed instructions on specific tasks
- **Conceptual Guides:** Dives into fundamental library concepts

These well-crafted, beginner-friendly documents are regularly updated to reflect the latest library changes.

The Docs main page displays a search bar and available documentation.

 **Hugging Face**

[Models](#) [Datasets](#) [Spaces](#) [Docs](#) [Solutions](#) [Pricing](#)

Documentations

- Hub**
Host Git-based models, datasets and Spaces on the Hugging Face Hub.
- Transformers**
State-of-the-art ML for Pytorch, TensorFlow, and JAX.
- Diffusers**
State-of-the-art diffusion models for image and audio generation in PyTorch.
- Hub Python Library**
Client library for the HF Hub: manage repositories from your Python runtime.
- Datasets**
Access and share datasets for computer vision, audio, and NLP tasks.
- Gradio**
Build machine learning demos and other web apps, in just a few lines of Python.
- Inference API**
Experiment with over 200k models easily using our free Inference API.
- Huggingface.js**
A collection of JS libraries to interact with Hugging Face, with TS types included.
- Transformers.js**
Community library to run pretrained models from Transformers in your browser.
- Inference Endpoints**
Easily deploy models to production on dedicated, fully managed infrastructure.
- PEFT**
Parameter efficient finetuning methods for large models
- Accelerate**
- Optimum**
- AWS Trainium & Inferentia**

1. In the top navigation, click **Docs**.
2. On the **Docs** main page, in the left pane, click **Hub**.

 **Hugging Face**

[Models](#) [Datasets](#) [Spaces](#) [Docs](#) [Solutions](#) [Pricing](#)

Documentations

- Hub**
Host Git-based models, datasets and Spaces on the Hugging Face Hub.
- Transformers**
State-of-the-art ML for Pytorch, TensorFlow, and JAX.
- Diffusers**
State-of-the-art diffusion models for image and audio generation in PyTorch.
- Hub Python Library**
Client library for the HF Hub: manage repositories from your Python runtime.
- Datasets**
Access and share datasets for computer vision, audio, and NLP tasks.
- Gradio**
Build machine learning demos and other web apps, in just a few lines of Python.
- Inference API**
Experiment with over 200k models easily using our free Inference API.
- Huggingface.js**
A collection of JS libraries to interact with Hugging Face, with TS types included.
- Transformers.js**
Community library to run pretrained models from Transformers in your browser.
- Inference Endpoints**
Easily deploy models to production on dedicated, fully managed infrastructure.
- PEFT**
Parameter efficient finetuning methods for large models
- Accelerate**
- Optimum**
- AWS Trainium & Inferentia**

3. The Hugging Face Hub is a central platform with over 120k models, 20k datasets, and 50k demo apps (Spaces) that are open-source and publicly available. On the Hugging Face Hub documentation page, you can view comprehensive documents organized into five sections: Repositories, Models, Datasets, Spaces, and Other. Users can easily collaborate and build ML together on the Hugging Face Hub.

Hub

Search documentation

EN

178

Hugging Face Hub

Repositories

Getting Started with Repositories

Repository Settings

Pull Requests & Discussions

Notifications

Collections

Webhooks

Repository size recommendations

Next Steps

Licenses

Models

The Model Hub

Model Cards

Gated Models

Uploading Models

Downloading Models

Integrated Libraries

Model Widgets

Inference API docs

Frequently Asked Questions

Advanced Topics

Hugging Face Hub documentation

The Hugging Face Hub is a platform with over 120k models, 20k datasets, and 50k demo apps (Spaces), all open source and publicly available, in an online platform where people can easily collaborate and build ML together. The Hub works as a place where anyone can explore, experiment, collaborate, and build technology with Machine Learning. Are you ready to start your path towards open source Machine Learning? 🤖

Repositories

Introduction

Getting Started

Repository Settings

Pull requests and Discussions

Notifications

Collections

Webhooks

Next Steps

Licenses

Models

Introduction

The Model Hub

Model Cards

Gated Models

Libraries

Tasks

Uploading Models

Downloading Models

Libraries

Tasks

Widgets

Inference API

Datasets

Introduction

Datasets Overview

Dataset Cards

Gated Datasets

Dataset viewer

Using Datasets

Adding New Datasets

Spaces

Introduction

Other

Organizations

Step 10: Try yourself

Read through the documents that interest you and get started on your ML journey!

Step 11: Exploring Solutions

Hugging Face provides businesses and individuals with a variety of solutions for developing and deploying ML models:

Enterprise Hub: A paid subscription service that offers enhanced features like single sign-on (SSO) and security assertion markup language (SAML) support, audit logs, customizable storage, and deployment on user infrastructure with managed billing


Expert Acceleration Program (EAP): A paid program to gain direct access to Hugging Face experts for accelerating ML projects

Inference Endpoints: A set of pre-trained models for NLP tasks that are easily deployable in production

AutoTrain: A no-code tool for training modern ML models with a focus on NLP

Hardware: Hugging Face collaborates with AI hardware accelerators to improve model training and deployment speed

1. In the top navigation, click **Solutions** to view the drop-down list with the five solutions: Enterprise Hub, Expert Acceleration Program, Inference Endpoints, AutoTrain, and Hardware.



Step 12: Exploring Pricing

Let's move to the last feature, Pricing.

Hugging Face offers a range of products and services, some free and others paid (subscription or pay-as-you-go).

- The **free services** include the Hugging Face Hub, Transformers library, Datasets, and Models.
- The **paid services** encompass Spaces Hardware, Inference Endpoints, AutoTrain, and Enterprise Hub.

The pricing for the paid services depends on resource usage, with varying rates for different hardware types and services.

Examples of Hugging Face pricing:

- **Spaces Hardware:** CPU basic machines start at \$0.05/hour; GPU machines start at \$0.20/hour, and accelerator machines start at \$1.50/hour.
- **Inference Endpoints:** CPU basic machines start at \$0.06/hour; GPU machines start at \$0.25/hour, and accelerator machines start at \$2.00/hour.
- **Enterprise Hub:** Starts at \$20/month per user.

 **Hugging Face**

Search models, datasets, users...

ModelsDatasetsSpacesDocsSolutionsPricing

Pricing

The simplest way to access compute for AI

Users and organizations already use the Hub as a collaboration platform, we're making it easy to seamlessly and scalably launch ML compute directly from the Hub.

**HF Hub**
Collaborate on Machine Learning

- ✓ Host unlimited models, datasets, and Spaces
- ✓ Create unlimited orgs and private repos
- ✓ Access the latest ML tools and open source
- ✓ Community support

Forever
Free

PRO Pro Account

Show your support for the best ML community

- ✓ Get a PRO badge on your profile
- ✓ Early access to new features
- ✓ Unlock Inference for PROs
- ✓ Higher tier for AutoTrain

Subscribe for
\$9 /month

**Enterprise Hub**
Accelerate your AI roadmap

- ✓ SSO and SAML support
- ✓ Audit logs, storage location (EU, US, Asia)
- ✓ Deploy Inference on your own Infra
- ✓ Managed billing with yearly commits

Starting at
\$20 per user per month

 **Spaces Hardware**

 **Inference Endpoints**

 **AutoTrain**

Summary

Congratulations on completing the hands-on lab, Explore Hugging Face.

In this lab, you explored the various features and capabilities of Hugging Face, the open-source AI community and data science platform. You also looked at the Hugging Face features: Models, Databases, Spaces, Docs, Solutions, and Pricing in detail.

© IBM Corporation. All rights reserved.



Skills Network